



# **Comp90042**

## **Natural Language Processing**

### **Workshop**

### **(Week3 | 2025s1)**

*Dr Jean Lee*



# Agenda

## 1. Discussion (~45 mins)

- **Text Classification**
- **N-gram Language Model**

## 2. Programming (~10 mins)

3	 <a href="#">workshop-03.pdf</a> ↓	<a href="#">03-classification.ipynb</a> ↓ <a href="#">04-ngram.ipynb</a> ↓
---	---	---

# Discussion : Text classification

1. What is **text classification**? Give some examples.
  - (a) Why is text classification generally a difficult problem? What are some hurdles that need to be overcome?
  - (b) Consider some (supervised) text classification problem, and discuss whether the following (supervised) machine learning models would be suitable:
    - i.  $k$ -Nearest Neighbour using Euclidean distance
    - ii.  $k$ -Nearest Neighbour using Cosine similarity
    - iii. Decision Trees using Information Gain
    - iv. Naive Bayes
    - v. Logistic Regression
    - vi. Support Vector Machines

# Text Classification

## What is Text Classification?

- Is the task of classifying text documents into different labels.

### Input


- a document  $d$
- a fixed set of classes (labels)  $C = \{c_1, c_2, \dots, c_j\}$

(ML run) a training set of  $m$  labeled documents  $(d_1, c_1), \dots, (d_m, c_m)$

### Output

- a predicted class  $c \in C$

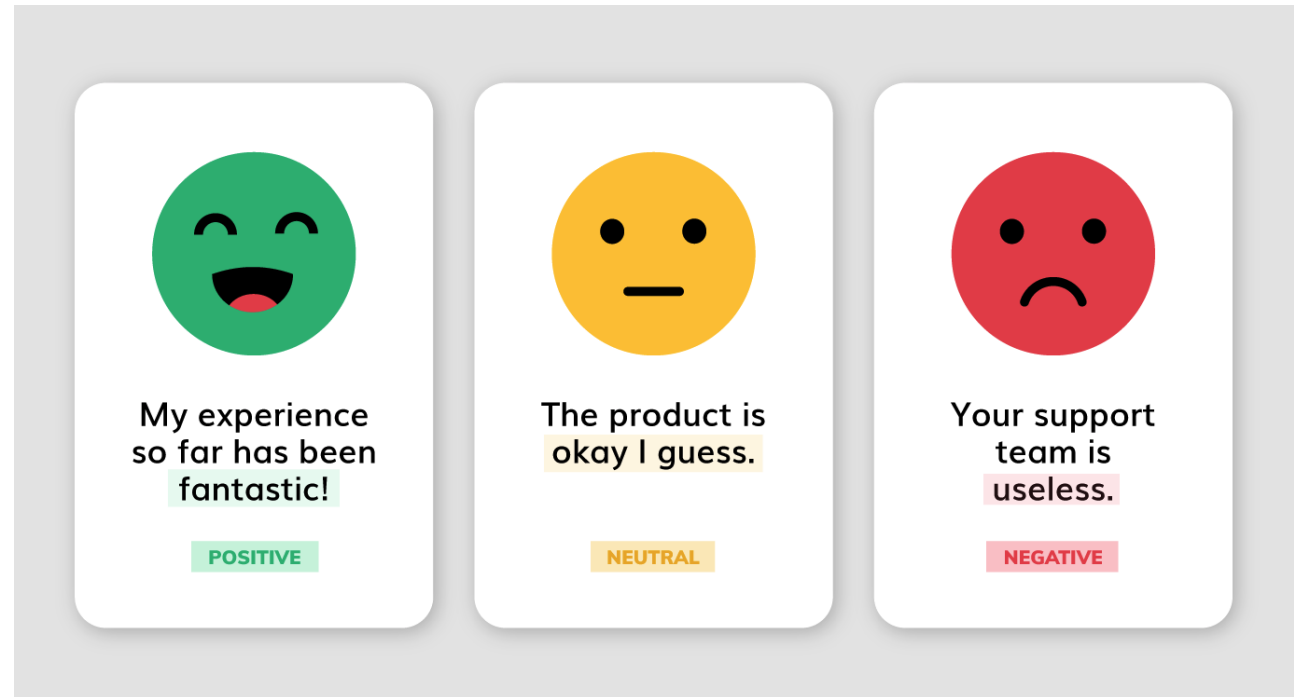
(ML run) a learned classifier  $\gamma: d \rightarrow c$



# Text Classification

## Text Classification Examples

- Sentiment analysis
- Spam detection
- Topic classification
- Authorship identification
- Native-language identification
- Automatic fact-checking
- ...



Three cards illustrating sentiment analysis results:

- Positive:** My experience so far has been fantastic! (Green smiley face icon, POSITIVE label)
- Neutral:** The product is okay I guess. (Yellow neutral face icon, NEUTRAL label)
- Negative:** Your support team is useless. (Red sad face icon, NEGATIVE label)

# Text Classification - Challenge

## Why is text classification generally a difficult problem?

- The main issue is in terms of document representation.
- *how do we identify **features** of the **document** which help us to distinguish between the various classes?*
- **Source** of document features: tokens (words) in the document
  - **feature selection** is often important
  - Single words: inadequate at modelling meaningful information
  - Multi-word (e.g. bi-grams): sparse data problem

# Text Classification – Word Representation

## How do we represent the meaning of a word?

- **Count-based word representation** : (e.g.) One-hot vectors, Bag of Words, Term Frequency-Inverse Document Frequency (TF-IDF)
- **Prediction-based word representation** : (e.g.) Word2Vec, GloVe

week5

## One-hot vectors (encoding)

- regard words as discrete symbols:

motel = [0 0 0 0 0 0 0 0 0 0 1 0 0 0 0]

hotel = [0 0 0 0 0 1 0 0 0 0 0 0 0 0 0]

### *issues:*

- ➔ *no natural notion of similarity*
- ➔ *number of words in vocabulary (inefficiency)*

# Text Classification – Word Representation

## How do we represent the meaning of a word?

### Bag of Words (BoW)

- a representation of text that describes **the occurrence of words** within a document.
- The intuition is that documents are similar if they have similar content.

S1= I **love** you but you **hate** me

S2= I **hate** you but you **love** me



#### *issues:*

- ➔ *Discarding word order*
- ➔ *ignores the context*
- ➔ *ignores meaning of words in the document (semantics).*





# Text Classification – ML models

## Which ML models would be suitable for text classification?

*Select one classifier and explain it to your partner. Each of you should choose a different classifier.*

- k-Nearest Neighbour (kNN) using Euclidean distance
- k-Nearest Neighbour (kNN) using Cosine similarity
- Decision Trees using Information Gain
- Naive Bayes
- Logistic Regression
- Support Vector Machines

### Prerequisites


- Machine learning basics (COMP30027, COMP90049, COMP90051)
  - Modules → Welcome → Machine Learning and Linguistics Readings

[https://canvas.lms.unimelb.edu.au/courses/210955/pages/machine-learning-and-linguistics-readings?module\\_item\\_id=6453877](https://canvas.lms.unimelb.edu.au/courses/210955/pages/machine-learning-and-linguistics-readings?module_item_id=6453877)

<https://www.youtube.com/watch?v=E0Hmnixke2g>

# k-Nearest Neighbour (kNN)

- KNN: Classify based on **majority class of k-nearest** training examples in feature space.




- Distance metric
- Choosing k (e.g. k = 3 or ?)

(+) *easy to implement, few hyperparameters*

(-) *not good for scale and high-dimensional data, overfitting*


- k-Nearest Neighbour using **Euclidean distance**
  - Often this is a **bad idea**
  - tends to classify documents **based upon their length**
  - which is usually not a distinguishing characteristic for text classification problems.



$$\text{Euclidean Distance (d)} = \sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2}$$

# k-Nearest Neighbour (kNN)

- k-Nearest Neighbour using **Cosine similarity**
  - Usually better than using Euclidean distance
  - However, **kNN** suffers from **high-dimensionality problems**
  - our feature set based upon the presence of (all) words usually *isn't suitable* for this model.
- **Cosine similarity**
  - measures the similarity between two vectors of an inner (dot) product space.
  - the cosine of the angle between two vectors.

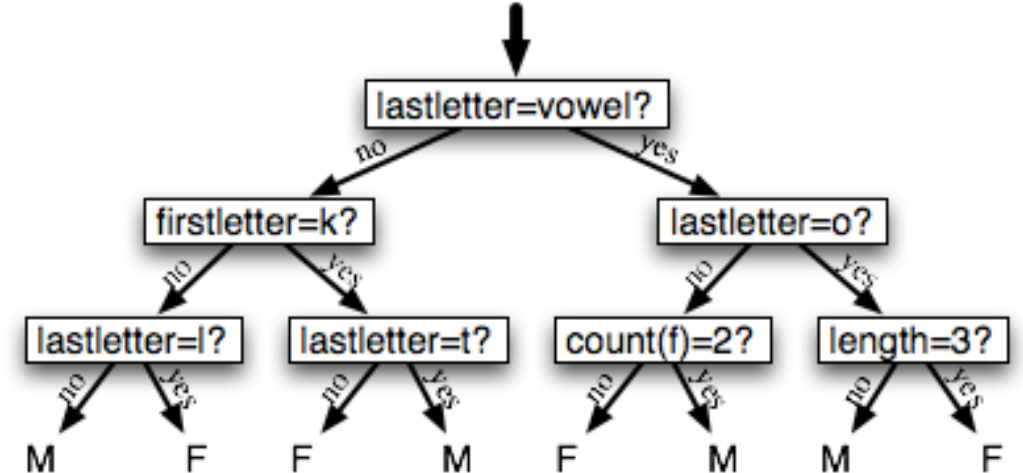


$$\cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}}$$

$$\begin{aligned} \cos(\text{Queen}, \text{King}) &= \frac{(0.3 \cdot 0.5) + (0.9 \cdot 0.7)}{\sqrt{0.3^2 + 0.9^2} \cdot \sqrt{0.5^2 + 0.7^2}} \\ &= \frac{0.15 + 0.63}{\sqrt{0.9^2} \cdot \sqrt{0.74}} \\ &= \frac{0.78}{\sqrt{0.666}} \\ &= 0.03 \end{aligned}$$


# Decision Trees (DT)

- DT : Construct a **tree** where **nodes** correspond to **tests on individual features**
  - **can be useful** for finding meaningful features
  - However, the feature set is very large, and we might find spurious correlations.
- Decision Trees using **Information Gain**
  - **Information Gain** : to determine the best features for splitting nodes.
  - **poor choice** because it tends to prefer **rare** features.



# Naive Bayes (NB)


- NB : Find the class with **the highest likelihood under Bayes Law**
  - At first glance, a **poor choice**
  - the "naive" assumption of the **conditional independence** of features and classes is highly **untrue**.
  - Also sensitive to a large feature set
  - Surprisingly somewhat **useful anyway!**



# Logistic Regression (LR)


- LR : Put linear combination of features in **logistic function**
  - **Useful**
  - it **relaxes the conditional independence requirement** of Naive Bayes.
  - Can handle large numbers of mostly useless features by **‘feature weighting’ step**

$$y = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x)}}$$



# Support Vector Machines (SVM)

- SVM : Finds **hyperplane** which **separates the training data with maximum margin**
  - Linear kernels often quite *effective*
  - Some combination of features are useful for characterising the classes.
  - Problems: **multiple classes** (most text classification tends to be multi-class).



# Discussion : N-gram Language Model (LM)

2. For the following “corpus” of two documents:

1. how much wood would a wood chuck chuck if a wood chuck would chuck wood
2. a wood chuck would chuck the wood he could chuck if a wood chuck would chuck wood

(a) Which of the following sentences: a wood could chuck; wood would a chuck; is more probable, according to:

- i. An unsmoothed uni-gram language model?
- ii. A uni-gram language model, with Laplacian (“add-one”) smoothing?
- iii. An unsmoothed bi-gram language model?
- iv. A bi-gram language model, with Laplacian smoothing?
- v. An unsmoothed tri-gram language model?
- vi. A tri-gram language model, with Laplacian smoothing?

(b) Assuming we are using a bi-gram language model with Kneser-Ney smoothing. Given the bigram `chuck a`, compute the continuation probability for `a`.

3. What does **back-off** mean, in the context of smoothing a language model? What does **interpolation** refer to?

	unsmooth	Laplacian
unigram	1)	2)
bi-gram	3)	4)
tri-gram	5)	6)



# N-gram LM Calculation

- Corpus

1: how much wood would a wood chuck chuck if a wood chuck would chuck wood

2: a wood chuck would chuck the wood he could chuck if a wood chuck would chuck wood

- Word counts

<s>	a	chuck	could	he	how	if	much	the	wood	would	</s>	Total
2	4	9	1	1	1	2	1	1	8	4	2	34

- M: the number of all words - sum of all word frequency ( $M = 34$ )
- V: the number of unique words - vocabulary size ( $V = 11$ )

## Why is the start symbol <s> left out?

- used internally to set context but not included in the final n-gram probabilities
- because they are **artificially inserted** and **do not represent actual text content**.
- While the end symbol </s> is typically included because </s> **helps the model learn the likelihood of ending** after certain sequences of words.

# N-gram: 1) Unigram unsmooth

a	chuck	could	he	how	if	much	the	wood	would	</s>	Total
4	9	1	1	1	2	1	1	8	4	2	34

- M: the number of all words - sum of all word frequency ( **$M = 34$** )
- V: the number of unique words - vocabulary size ( **$V = 11$** )

- Unigram probability (Unsmooth)

$$P(w_i) = \frac{C(w_i)}{M}$$

**$M$**  ← Total number of word tokens in corpus

**A: a wood could chuck**

**B: wood would a chuck**

$$\begin{aligned} P(A) &= P(a)P(\text{wood})P(\text{could})P(\text{chuck})P(</s>) \\ &= \frac{4}{34} \times \frac{8}{34} \times \frac{1}{34} \times \frac{9}{34} \times \frac{2}{34} \approx 1.27 \times 10^{-5} \end{aligned}$$

$$\begin{aligned} P(B) &= P(\text{wood})P(\text{would})P(a)P(\text{chuck})P(</s>) \\ &= \frac{8}{34} \times \frac{4}{34} \times \frac{4}{34} \times \frac{9}{34} \times \frac{2}{34} \approx 5.07 \times 10^{-5} \end{aligned}$$

# N-gram: 2) Unigram Laplacian smoothing

a	chuck	could	he	how	if	much	the	wood	would	</s>	Total
4	9	1	1	1	2	1	1	8	4	2	34

- M: the number of all words - sum of all word frequency ( **$M = 34$** )
- V: the number of unique words - vocabulary size ( **$V = 11$** )

- Unigram Laplacian (“add-one”) smoothing

$$P_{add1}(w_i) = \frac{C(w_i) + 1}{M + |V|}$$

**A: a wood could chuck**

**B: wood would a chuck**

$$\begin{aligned} P_L(A) &= P_L(a)P_L(\text{wood})P_L(\text{could})P_L(\text{chuck})P_L(</s>) \\ &= \frac{5}{45} \times \frac{9}{45} \times \frac{2}{45} \times \frac{10}{45} \times \frac{3}{45} \approx 1.46 \times 10^{-5} \end{aligned}$$

$$\begin{aligned} P_L(B) &= P_L(\text{wood})P_L(\text{would})P_L(a)P_L(\text{chuck})P_L(</s>) \\ &= \frac{9}{45} \times \frac{5}{45} \times \frac{5}{45} \times \frac{10}{45} \times \frac{3}{45} \approx 3.66 \times 10^{-5} \end{aligned}$$

# N-gram: 3) bi-gram unsmooth

1: how much wood would a wood chuck chuck if a wood chuck would chuck wood

2: a wood chuck would chuck the wood he could chuck if a wood chuck would chuck wood

a	chuck	could	he	how	if	much	the	wood	would	</s>	Total
4	9	1	1	1	2	1	1	8	4	2	34

- bi-gram probability (Unsmooth)

$$P(w_i | w_{i-1}) = \frac{C(w_{i-1}w_i)}{C(w_{i-1})}$$

A: a wood could chuck

B: wood would a chuck

$$\begin{aligned}
 P(A) &= P(a|<s>)P(\text{wood}|a)P(\text{could}|\text{wood})P(\text{chuck}|\text{could})P(</s>|\text{chuck}) \\
 &= \frac{1}{2} \times \frac{4}{4} \times \frac{0}{8} \times \frac{1}{1} \times \frac{0}{9} = 0 \\
 P(B) &= P(\text{wood}|<s>)P(\text{would}|\text{wood})P(a|\text{would})P(\text{chuck}|a)P(</s>|\text{chuck}) \\
 &= \frac{0}{2} \times \frac{1}{8} \times \frac{1}{4} \times \frac{0}{4} \times \frac{0}{9} = 0
 \end{aligned}$$

# N-gram: 4) bi-gram Laplacian smoothing

- 1: how much wood would a wood chuck chuck if a wood chuck would chuck wood
- 2: a wood chuck would chuck the wood he could chuck if a wood chuck would chuck wood

a	chuck	could	he	how	if	much	the	wood	would	</s>	Total
4	9	1	1	1	2	1	1	8	4	2	34

- M: the number of all words - sum of all word frequency ( $M = 34$ )
- V: the number of unique words - vocabulary size ( $V = 11$ )
- bi-gram Laplacian (“add-one”) smoothing

$$P_{add1}(w_i | w_{i-1}) = \frac{C(w_{i-1}w_i) + 1}{C(w_{i-1}) + |V|}$$

A: a wood could chuck

B: wood would a chuck

$$\begin{aligned}
 P_L(A) &= P_L(a|<s>)P_L(\text{wood}|a)P_L(\text{could}|wood)P_L(\text{chuck}|could) \\
 &\quad P_L(</s>|chuck) \\
 &= \frac{2}{13} \times \frac{5}{15} \times \frac{1}{19} \times \frac{2}{12} \times \frac{1}{20} \approx 2.25 \times 10^{-5} \\
 P_L(B) &= P_L(\text{wood}|<s>)P_L(\text{would}|wood)P_L(a|would)P_L(\text{chuck}|a) \\
 &\quad P_L(</s>|chuck) \\
 &= \frac{1}{13} \times \frac{2}{19} \times \frac{2}{15} \times \frac{1}{15} \times \frac{1}{20} \approx 3.60 \times 10^{-6}
 \end{aligned}$$

# N-gram: 5) tri-gram unsmooth

a	chuck	could	he	how	if	much	the	wood	would	</s>	Total
4	9	1	1	1	2	1	1	8	4	2	34

- M: the number of all words - sum of all word frequency ( **$M = 34$** )
- V: the number of unique words - vocabulary size ( **$V = 11$** )

## • tri-gram probability (Unsmooth)

$$P(w_i | w_{i-1} w_{i-2}) = \frac{C(w_{i-2} w_{i-1} w_i)}{C(w_{i-2} w_{i-1})}$$

**A: a wood could chuck**

**B: wood would a chuck**

$$\begin{aligned} P(A) &= P(a | <s> <s>) P(\text{wood} | <s> a) \cdots P(</s> | \text{could chuck}) \\ &= \frac{1}{2} \times \frac{1}{1} \times \frac{0}{4} \times \frac{0}{0} \times \frac{0}{1} = ? \end{aligned}$$

$$\begin{aligned} P(B) &= P(\text{wood} | <s> <s>) P(\text{would} | <s> \text{wood}) \cdots P(</s> | a \text{ chuck}) \\ &= \frac{0}{2} \times \frac{0}{0} \times \frac{1}{1} \times \frac{0}{1} \times \frac{0}{0} = ? \end{aligned}$$

# N-gram: 6) tri-gram Laplacian smoothing

a	chuck	could	he	how	if	much	the	wood	would	</s>	Total
4	9	1	1	1	2	1	1	8	4	2	34

- M: the number of all words - sum of all word frequency ( **$M = 34$** )
- V: the number of unique words - vocabulary size ( **$V = 11$** )

## • tri-gram Laplacian (“add-one”) smoothing

$$P_L(w_i | w_{i-1} w_{i-2}) = \frac{C(w_{i-2} w_{i-1} w_i) + 1}{C(w_{i-2} w_{i-1}) + V}$$

**A: a wood could chuck**

**B: wood would a chuck**

$$\begin{aligned} P_L(A) &= P_L(a | <s> <s>) P_L(\text{wood} | <s> a) \cdots P_L(</s> | \text{could chuck}) \\ &= \frac{2}{13} \times \frac{2}{12} \times \frac{1}{15} \times \frac{1}{11} \times \frac{1}{12} \approx 1.30 \times 10^{-5} \end{aligned}$$

$$\begin{aligned} P_L(B) &= P_L(\text{wood} | <s> <s>) P_L(\text{would} | <s> \text{wood}) \cdots P_L(</s> | a \text{ chuck}) \\ &= \frac{1}{13} \times \frac{1}{11} \times \frac{2}{12} \times \frac{1}{12} \times \frac{1}{11} \approx 8.83 \times 10^{-6} \end{aligned}$$

# N-gram: bi-gram Kneser-Ney (KN) smoothing

- Given the bigram chuck a, compute the continuation probability for a

1: how much wood would a wood chuck chuck if a wood chuck would chuck wood

2: a wood chuck would chuck the wood he could chuck if a wood chuck would chuck wood

a	chuck	could	he	how	if	much	the	wood	would	</s>	Total
4	9	1	1	1	2	1	1	8	4	2	34

- Observed bi-grams:** chuck chuck, chuck if, chuck would, chuck wood, chuck the
- Unobserved bi-grams:** chuck a, chuck could, chuck he, chuck how, chuck much, chuck </s>

- Continuation counts** *number of unique words in the vocabulary which appears before a word w*

- a = {would, if, <s>} = 3
- could = {he} = 1
- he = {wood} = 1

- how = {<s>} = 1
- much = {how} = 1
- </s> = {wood} = 1

$$P_{cont}(a) = \frac{\#_{cont}(a)}{\#_{cont}(a) + \#_{cont}(could) + \#_{cont}(he) + \#_{cont}(how) + \#_{cont}(much) + \#_{cont}(</s>)}$$

$$= \frac{3}{3 + 1 + 1 + 1 + 1 + 1}$$

**Numerator:** continuation count of ( $w_i$ )

**Denominator:** Sum of the continuation count of all unobserved word for ( $w_{i-1}$ )



# N-gram: bi-gram Kneser-Ney (KN) smoothing

- bi-gram probability (Unsmooth)

$$P(w_i | w_{i-1}) = \frac{C(w_{i-1}w_i)}{C(w_{i-1})}$$

- bi-gram Laplacian (“add-one”) smoothing

$$P_{add1}(w_i | w_{i-1}) = \frac{C(w_{i-1}w_i) + 1}{C(w_{i-1}) + |V|}$$

- bi-gram Kneser-Ney (KN) smoothing

$$P_{KN}(w_i | w_{i-1}) = \begin{cases} \frac{C(w_{i-1}, w_i) - D}{C(w_{i-1})}, & \text{if } C(w_{i-1}, w_i) > 0 \\ \beta(w_{i-1}) P_{cont}(w_i), & \text{otherwise} \end{cases}$$

the amount of probability mass that  
has been discounted for context  $w_{i-1}$

$$P_{cont}(w_i) = \frac{|\{w'_{i-1} : C(w'_{i-1}, w_i) > 0\}|}{\sum_{\{w_j : C(w_{i-1}, w_j) = 0\}} |\{w_{j-1} : C(w_{j-1}, w_j) > 0\}|}$$

# N-gram: back-off and interpolation

## Back-off

- Use **lower-order n-gram model** if higher-order is unseen.
- For example, if we have never seen some tri-gram from our sentence, we can instead consider the bigram probability.

## interpolation

- Take **weighted average sum of n-gram**.
- Instead of only “falling back” to lower (back-off), consider every probability as a linear combination of all of the relevant n-gram models.

$$\begin{aligned} P_{\text{Interpolation}}(w_m \mid w_{m-1}, w_{m-2}) &= \lambda_3 p_3^*(w_m \mid w_{m-1}, w_{m-2}) \\ &\quad + \lambda_2 p_2^*(w_m \mid w_{m-1}) \\ &\quad + \lambda_1 p_1^*(w_m). \end{aligned}$$

# Programming!

## Programming

1. In the `03-classification` notebook, observe how different tokenisation regimes alter the text classification performance of the various classifiers on the given Reuters dataset problem.
2. Using the iPython notebook `04-ngram`, randomly generate some sentences based on the bi-gram models of the Gutenberg corpus and the Penn Treebank. What do you notice about these sentences? Are there any sentences which might get returned for both corpora? Why?